

Docker Layers Explained

by Gunter Rotsaert  MVB · Mar. 18, 19 · Cloud Zone · Tutorial

Discover the future of enterprise automation in the hybrid multi-cloud era. Create the connected, secure, and flexible application environment you need. [Download](#) the 451 Research brief to learn how.

Presented by Red Hat

When you pull a Docker image, you will notice that it is pulled as different layers. Also, when you create your own Docker image, several layers are created. In this post we will try to get a better understanding of Docker layers.

1. What Is a Docker Layer?

A Docker image consists of several layers. Each layer corresponds to certain instructions in your `Dockerfile`. The following instructions create a layer: `RUN`, `COPY`, `ADD`. The other instructions will create intermediate layers and do not influence the size of your image. Let's take a look at an example. We will use a Spring Boot MVC application which we have created beforehand and where the Maven build creates our Docker image. The sources are available at [GitHub](#).

We use the `feature/dockerbenchsecurity` branch, which is a more secure version of the `master` branch. Here is the `Dockerfile`:

```
1 FROM openjdk:10-jdk
2 VOLUME /tmp
3
4 RUN useradd -d /home/appuser -m -s /bin/bash appuser
5 USER appuser
6
7 HEALTHCHECK --interval=5m --timeout=3s CMD curl -f http://localhost:8080/actuator/health/ ||
8
9 ARG JAR_FILE
10 COPY ${JAR_FILE} app.jar
11 ENTRYPOINT ["java","-Djava.security.egd=file:/dev/./urandom","-jar","/app.jar"]
```

We build the application with `mvn clean install` which will also create the Docker image. We don't list the pulling of all layers of the `openjdk:10-jdk` image for brevity.

```
1  Image will be built as mydeveloperplanet/mykubernetesplanet:0.0.3-SNAPSHOT
2
3  Step 1/8 : FROM openjdk:10-jdk
4
5  Pulling from library/openjdk
6  Image 16e82e17faef: Pulling fs layer
7  ...
8  Image a9448aba0bc3: Pull complete
9  Digest: sha256:9f17c917630d5e95667840029487b6561b752f1be6a3c4a90c4716907c1aad65
10 Status: Downloaded newer image for openjdk:10-jdk
11 ---> b11e88dd885d
12 Step 2/8 : VOLUME /tmp
13
14 ---> Running in 21329898c3a6
15 Removing intermediate container 21329898c3a6
16 ---> b6f9ca000de6
17 Step 3/8 : RUN useradd -d /home/appuser -m -s /bin/bash appuser
18
19 ---> Running in 82645047e6e7
20 Removing intermediate container 82645047e6e7
21 ---> 04f6b2716819
22 Step 4/8 : USER appuser
23
24 ---> Running in 697b663dadbb
25 Removing intermediate container 697b663dadbb
26 ---> eaf6b8af5709
27 Step 5/8 : HEALTHCHECK --interval=5m --timeout=3s CMD curl -f http://localhost:8080/actuator,
28
29 ---> Running in f420b9d060c5
30 Removing intermediate container f420b9d060c5
31 ---> 77f95436a3ff
32 Step 6/8 : ARG JAR_FILE
33
34 ---> Running in 60b9d25ad2ac
35 Removing intermediate container 60b9d25ad2ac
36 ---> 135fa7df95ac
37 Step 7/8 : COPY ${JAR_FILE} app.jar
38
39 ---> 63c18567012b
Step 8/8 : ENTRYPOINT ["java","-Djava.security.egd=file:/dev/./urandom","-jar","/app.jar"]
```

```

40
41
42 ---> Running in 79203446934a
43 Removing intermediate container 79203446934a
44 ---> 8e2b049f9783
45 Successfully built 8e2b049f9783
46 Successfully tagged mydeveloperplanet/mykubernetesplanet:0.0.3-SNAPSHOT

```

What happens here? We notice that layers are created and most of them are removed (*removing intermediate container*). So why does it say *removing intermediate container* and not *removing intermediate layer*? That's because a build step is executed in an intermediate container. When the build step is finished executing, the intermediate container can be removed. Besides that, a layer is read-only. A layer contains the differences between the preceding layer and the current layer. On top of the layers, there is a writable layer (the current one) which is called the *container layer*. As mentioned before, only specific instructions will create a new layer. Let's take a look at our Docker images:

```

1 $ docker image ls

```

REPOSITORY	TAG	IMAGE ID	CREATED
mydeveloperplanet/mykubernetesplanet	0.0.3-SNAPSHOT	8e2b049f9783	About a minute ago
openjdk	10-jdk	b11e88dd885d	2 months ago

And take a look at the history of our *mykubernetesplanet* Docker image:

```

1 $ docker history 8e2b049f9783

```

IMAGE	CREATED	CREATED BY	SIZE
8e2b049f9783	About a minute ago	/bin/sh -c #(nop) ENTRYPOINT ["java" "-Djav...	0B
63c18567012b	About a minute ago	/bin/sh -c #(nop) COPY file:2a5b71774c60e0f6...	17.4MB
135fa7df95ac	About a minute ago	/bin/sh -c #(nop) ARG JAR_FILE	0B
77f95436a3ff	2 minutes ago	/bin/sh -c #(nop) HEALTHCHECK &{["CMD-SHELL...	0B
eaf6b8af5709	2 minutes ago	/bin/sh -c #(nop) USER appuser	0B
04f6b2716819	2 minutes ago	/bin/sh -c useradd -d /home/appuser -m -s /b...	399kB
b6f9ca000de6	2 minutes ago	/bin/sh -c #(nop) VOLUME [/tmp]	0B
b11e88dd885d	2 months ago	/bin/sh -c #(nop) CMD ["jshell"]	0B
<missing>	2 months ago	/bin/sh -c set -ex; if [! -d /usr/share/m...	697MB

```
12 <missing>          2 months ago          /bin/sh -c #(nop) ENV JAVA_DEBIAN_VERSION=1...    0B
13 ...
```

We notice here, that the intermediate containers do have a size of 0B, just what was expected. Only the RUN and COPY command from the Dockerfile contribute to the size of the Docker image. The layers of the openjdk:10-jdk image are also listed and are recognized by the *missing* keyword. This only means that those layers are built on a different system and are not available locally.

2. Recreate the Docker Image

What happens if we run our Maven build again without making any changes to the sources?

```
1  Image will be built as mydeveloperplanet/mykubernetesplanet:0.0.3-SNAPSHOT
2
3  Step 1/8 : FROM openjdk:10-jdk
4
5  Pulling from library/openjdk
6  Digest: sha256:9f17c917630d5e95667840029487b6561b752f1be6a3c4a90c4716907c1aad65
7  Status: Image is up to date for openjdk:10-jdk
8  ---> b11e88dd885d
9  Step 2/8 : VOLUME /tmp
10
11 ---> Using cache
12 ---> b6f9ca000de6
13 ...
14 Step 6/8 : ARG JAR_FILE
15
16 ---> Using cache
17 ---> 135fa7df95ac
18 Step 7/8 : COPY ${JAR_FILE} app.jar
19
20 ---> 409f2fee0cde
21 Step 8/8 : ENTRYPOINT ["java","-Djava.security.egd=file:/dev/./urandom","-jar","/app.jar"]
22
23 ---> Running in 75f07955bbc8
24 Removing intermediate container 75f07955bbc8
25 ---> e5d7b72aad05
26 Successfully built e5d7b72aad05
27 Successfully tagged mydeveloperplanet/mykubernetesplanet:0.0.3-SNAPSHOT
```

We notice that the first layers are identical to our previous build. The layer ID's are the same. In the log, we

notice that the layers are taken from the cache. In step 7 a new layer is created with a new ID. We did create a new JAR file, Docker interprets this as a new file and therefore a new layer is created. In step 8 also a new layer is created, because it is build on top of the new layer.

Let's list the Docker images again:

```
1 $ docker image ls
```

REPOSITORY	TAG	IMAGE ID	CREATED	SIZE
mydeveloperplanet/mykubernetesplanet	0.0.3-SNAPSHOT	e5d7b72aad05	13 seconds ago	17.4MB
<none>	<none>	8e2b049f9783	5 minutes ago	0B
openjdk	10-jdk	b11e88dd885d	2 months ago	17.4MB

Our tag `0.0.3-SNAPSHOT` has received the *image ID* of our last build. The repository and tag of our old image ID are removed, which is indicated with the *none* keyword. This is called a *dangling* image. We will explain this in more detail at the end of this post.

When we take a look at the history of the newly-created image, we notice that the two top layers are new, just as in the build log:

```
1 $ docker history e5d7b72aad05
```

IMAGE	CREATED	CREATED BY	SIZE	COMMENT
e5d7b72aad05	38 seconds ago	/bin/sh -c #(nop) ENTRYPOINT ["java" "-Djav...	0B	
409f2fee0cde	42 seconds ago	/bin/sh -c #(nop) COPY file:4b04c6500d340c9e...	17.4MB	
135fa7df95ac	6 minutes ago	/bin/sh -c #(nop) ARG JAR_FILE	0B	
...				

When we make a source code change, the result is identical because also, in this case, a new JAR file is generated.

```
1 $ docker image ls
```

REPOSITORY	TAG	IMAGE ID	CREATED	SIZE
mydeveloperplanet/mykubernetesplanet	0.0.3-SNAPSHOT	eced642d4f5c	30 seconds ago	17.4MB
<none>	<none>	e5d7b72aad05	3 minutes ago	17.4MB
<none>	<none>	8e2b049f9783	8 minutes ago	0B
openjdk	10-jdk	b11e88dd885d	2 months ago	17.4MB

```

6
7 $ docker history eced642d4f5c
8
9 IMAGE                CREATED              CREATED BY          SIZE
10 eced642d4f5c         About a minute ago  /bin/sh -c #(nop) ENTRYPOINT ["java" "-Djav...  0B
11 44a9097b8bad         About a minute ago  /bin/sh -c #(nop) COPY file:1d5276778b53310e... 17.4MB
12 135fa7df95ac         9 minutes ago      /bin/sh -c #(nop) ARG JAR_FILE                0B
13 ...

```

3. What About Size?

Let's take a closer look at the latest output of the `docker image ls` command. We notice two dangling images, which are 1 GB in size. But what does this really mean for storage? First, we need to know where the data for our images are stored. Use the following command in order to retrieve the storage location:

```

1 $ docker image inspect eced642d4f5c
2 ...
3 "GraphDriver": {
4   "Data": {
5     "LowerDir": "/var/lib/docker/overlay2/655be8bea8e54c31ebb7e3adf05db227d194a49c1e2f9555:
6     "MergedDir": "/var/lib/docker/overlay2/205b55ee2f0e06394b6d17067338845410609887ccd18f53:
7     "UpperDir": "/var/lib/docker/overlay2/205b55ee2f0e06394b6d17067338845410609887ccd18f53:
8     "WorkDir": "/var/lib/docker/overlay2/205b55ee2f0e06394b6d17067338845410609887ccd18f53:
9   },
10   "Name": "overlay2"
11 },
12 ...

```

Our Docker images are stored at `/var/lib/docker/overlay2`. We can simply retrieve the size of the `overlay2` directory in order to have an idea of the allocated storage:

```

1 $ du -sh -m overlay2
2 1059 overlay2

```

The `openjdk:10-jdk` image is 987 MB. The JAR file in our image is 17.4 MB. The total size should be somewhere around 987 MB + 3 * 17.4 MB (two dangling images and one real image). This is about 1,040 MB. We can conclude that there is some sort of smart storage for the Docker images and that we cannot

simply add all the sizes of the Docker images in order to retrieve the real storage size. The difference is due to the existence of intermediate images. These can be revealed as follows:

```
1 $ docker images -a
```

REPOSITORY	TAG	IMAGE ID	CREATED	SIZE
mydeveloperplanet/mykubernetesplanet	0.0.3-SNAPSHOT	eced642d4f5c	7 days ago	1GB
<none>	<none>	44a9097b8bad	7 days ago	1GB
<none>	<none>	e5d7b72aad05	7 days ago	1GB
<none>	<none>	409f2fee0cde	7 days ago	1GB
<none>	<none>	8e2b049f9783	7 days ago	1GB
<none>	<none>	63c18567012b	7 days ago	1GB
<none>	<none>	135fa7df95ac	7 days ago	987Mi
<none>	<none>	77f95436a3ff	7 days ago	987Mi
<none>	<none>	eaf6b8af5709	7 days ago	987Mi
<none>	<none>	04f6b2716819	7 days ago	987Mi
<none>	<none>	b6f9ca000de6	7 days ago	987Mi
openjdk	10-jdk	b11e88dd885d	2 months ago	987Mi

4. Get Rid of Dangling Images

How do we get rid of these dangling images? We don't need them anymore and they only allocate storage. First, list the dangling images:

```
1 $ docker images -f dangling=true
```

REPOSITORY	TAG	IMAGE ID	CREATED	SIZE
<none>	<none>	e5d7b72aad05	7 days ago	1GB
<none>	<none>	8e2b049f9783	7 days ago	1GB

We can use the `docker rmi` command to remove the images:

```
1 $ docker rmi e5d7b72aad05
```

```
2 Deleted: sha256:e5d7b72aad054100d142d99467c218062a2ef3bc2a0994fb589f9fc7ff004afe
3 Deleted: sha256:409f2fee0cde9b5144f8e92887b61e49f3ccbd2b0e601f536941d3b9be32ff47
4 Deleted: sha256:2162a2af22ee26f7ac9bd95c39818312dc9714b8fbfbefb892ff827be15c7795b
```

Or you can use the `docker image prune` command to do so.

Now that we have removed the dangling images, let's take a look at the size of the `overlay2` directory:

```
1 $ du -sh -m overlay2
2 1026 overlay2
```

We got rid of 33 MB. This seems like it's not so much, but when you often build Docker images, this can grow significantly over time.


5. Conclusion

In this post we tried to get a better understanding of Docker layers. We noticed that intermediate layers are created and that dangling images remain on our system if we don't clean them regularly. We also inspected the size of the Docker images on our system.

Kubernetes Monitoring [Best Practices](#) - Ensure a reliable and highly efficient Kubernetes de production.

Presented by Svrdin

Like This Article? Read More From DZone



Understanding and Creating Effective Docker Images



Bleeding Edge on Docker




Docker image_build Module: Features for the Power User



**Free DZone Refcard
Cloud Native Data Grids: Hazelcast
IMDG With Kubernetes**

Topics: DOCKER BUILD , DOCKER , LAYERS , DANGLING IMAGES , INTERMEDIATE LAYERS , CLOUD

Published at DZone with permission of Gunter Rotsaert , DZone MVB. [See the original article here.](#) 
Opinions expressed by DZone contributors are their own.

Deployment Manifests That Work in